

Animint: a Grammar for Interactive Animations

Toby Dylan Hocking, Carson Sievert, Susan VanderPlas

Abstract—Animint is a new domain-specific language (DSL) for linked, interactive, animated plots. It builds on top of previous work on the grammar of graphics (ggplot2 in R) and interactive visualization using web standards (D3 in JavaScript). In other data visualization systems, interactive animations are difficult to accomplish, since they must be defined in terms of low-level operations and sometimes 100s of lines of code. In contrast, Animint’s high-level DSL can be used to produce a wide variety of complex interactive visualizations with only 10s of lines of code. It works by adding 2 new aesthetics to the grammar of graphics: `clickSelects`, which allows users to select elements of the plot, and `showSelected`, which displays only elements corresponding to the current selection. After using the declarative Animint DSL to define an interactive animation in R code, it is first compiled and then rendered in a web browser using D3. We discuss the design of Animint, then compare to related libraries, and show several example visualizations of high-dimensional time series data. A free/open-source implementation of Animint is available at <https://github.com/tdhock/animint>.

Index Terms—Information visualization, user interfaces, toolkits, 2D graphics, interactive, animation

1 INTRODUCTION

Interactive, animated data visualization is a useful tool for obtaining an intuitive understanding of patterns in multivariate data sets. In this paper, we introduce Animint, a new domain-specific language (DSL) for designing linked, interactive, animated data visualizations. To explain exactly what kinds of visualizations we are interested in creating, we begin with an illustrative example.

We are concerned with tabular data sets, with columns for each variable and rows for each observation or data point. Consider the World Bank data set, which consist of several economic variables measured for 205 countries from 1960 to 2012 [World Bank, 2012]. The visualization in Figure 1 shows six variables from this data set: year, country, life expectancy, fertility rate, region, and population. We have chosen to emphasize the year 1996 and the countries United States and Vietnam. The benefit of the interactive visualization is that the reader or user may interact with the plot to view data about another country or time period. Finally, since one of the variables in this visualization is the year of measurement, it is natural to use animation to show how the data changes with respect to the year variable. The purpose of Animint is to make it easy to design such interactive animations.

In this paper we will discuss three important roles that contribute to the creation of an Animint visualization: the developer, who implements the Animint library; the designer, who uses the Animint library to define a visualization; and the user, who selects data subsets to view in a web browser. The main goal of Animint is to provide an expressive language for designers, while allowing users the freedom to interact with the plot to selectively view data subsets of interest. The designer specifies data sets and maps variables to interactive visual elements using the Animint DSL, then uses the Animint library to compile and save an interactive animation. The user writes no code, but can view and interact with an Animint visualization by clicking Scalable Vector Graphics

(SVG) elements in a web browser. A final goal is that the Animint library developer should be the only role concerned with low-level plotting details such as drawing axes and legends, freeing the designer and user to focus on interpreting the data.

The Animint DSL is for visualizations which are both interactive and animated. For the purposes of this paper, we define “interactive” to mean that graphical elements such as data points may be clicked to update the data that is shown in related plots. This is closely related to the concept of “direct manipulation,” which was introduced by Shneiderman [1982] and later applied to statistical data visualization [Hutchins et al., 1985, Becker and Cleveland, 1987, Cleveland, 1993]. Our definition of an “animated” visualization is one that can be watched like a video, automatically updating over time.

The type of interactivity that we propose is closely related to the system described by Cleveland [1993]. In that system, the user’s mouse creates a single rectangular brush which either defines the selection (transient mode), or defines data points to add to or remove from the selection (lasting mode or erasing mode). Cleveland also defines several operations on the selection, such as labeling, deleting, and enhanced linking. For example, in a scatter plot matrix, Cleveland uses enhanced linking to highlight which points are selected in each plot.

In Animint, the central concept of interactivity is a selection variable, such as year or country in Figure 1. For each selection variable, one or several values can be selected at a time, e.g. `year=1996` and `country={United States, Vietnam}`. Like Cleveland’s system, Animint supports enhanced linking to highlight the selected value(s) of each selection variable. In contrast to Cleveland’s single rectangular brush that selects points in plots of a single data table, an Animint designer may designate several selection variables in plots of several linked data tables. To declare a clickable plot element that changes the selected value of the year variable, an Animint designer writes `clickSelects=year`.

To achieve data-driven animations, Animint defines one other important operation involving selection variables: showing and hiding subsets of data. For example, in the scatterplot of Figure 1 we draw a point for each country, and want to move and change the size of each point as the selected year changes. To accomplish this,

• Toby Dylan Hocking is with the McGill University department of Human Genetics, Susan VanderPlas and Carson Sievert are with Iowa State University department of Statistics.
E-mail: toby.hocking@mail.mcgill.ca.

an Animint designer can simply declare `showSelected=year` which means to show only the data with the selected value of year.

Using just the `clickSelects` and `showSelected` keywords, a wide variety of interactive visualizations can be defined. To make the selection automatically change over time (animation), an Animint designer may declare one variable as the time variable. Since it is perceptually advantageous to have smooth transitions in data-driven animations [Heer and Robertson, 2007], an Animint designer may also declare a `duration` list of selection variables which should have smooth transitions.

The Animint DSL is implemented as an extension of `ggplot2` [Wickham, 2009, 2010], which is an R implementation of the grammar of graphics [Wilkinson et al., 2006]. The `ggplot2` language was created as a high-level abstraction for non-interactive visualizations. One of the key strengths of `ggplot2` is that it allows a designer to declare a visualization in terms of distinct geometric elements, each with a clear aesthetic mapping from data variables to geometric properties. The Animint DSL extends `ggplot2` by adding `clickSelects` and `showSelected` aesthetics.

The Animint library includes a compiler that converts a list of `ggplots` to an interactive web visualization rendered using the Data-Driven Documents (D3) library for JavaScript [Bostock et al., 2011]. One of the main reasons for the success and popularity of D3 is that it allows visualizations to be specified using the terminology of the Document Object Model (DOM), which makes learning D3 easy for web designers. Animint’s DSL abandons the DOM standard and sacrifices some of the flexibility of D3, but makes it easier to express one kind of interaction: showing and hiding subsets of data. The Animint DSL thus enables design of a large class of complex interactive data visualizations using just a few lines of code.

Finally, another key strength of `ggplot2` and D3 for visualization design are the libraries’ declarative syntax. Heer and Bostock [2010] proposed a declarative syntax for animated transitions, and studied the benefits of declarative languages for data visualization.

Animint is another declarative DSL, but defined at a higher level of abstraction than D3. It enables designers to focus on data visualization, while the Animint library developers can work on improving the lower-level rendering details.

The rest of this paper is organized as follows: we discuss related work in Section 2, and the design of the Animint system in Section 3. Then we perform a detailed comparison with other R packages in Section 4, and show some example applications of Animint in Section 5. Finally, we share some user feedback in Section 6 and then discuss future work in Section 7.

2 RELATED WORK

In this section we offer a comparison between Animint and several related systems, focusing on libraries with free/open-source software implementations in JavaScript and R. The main difference is that Animint’s declarative DSL allows interactive animations to be expressed much more easily than the other libraries.

2.1 Types of selection

TODO For example, it is popular to select regions with brushing tools—see [Heer et al., 2008], [Willett et al., 2007] for lots of good examples of selection modalities. Which of these does the system support? Do the authors regret losing the ability to multi-select, or is the grammar theoretically expandable to adding those too? Are these unnecessary, or uninteresting? Are they HARD to express, or simply things that the authors didn’t get around to doing?

2.2 D3 and other JavaScript libraries

Animint uses D3 to render interactive animations in a web browser [Bostock et al., 2011]. D3 uses a lower level of abstraction than Animint, so is able to express a wider variety of visualizations. However, Animint can be used to more succinctly declare certain types of data visualizations. For example, Figure 2 shows some Animint and D3 code required to define the scatterplot in the




Fig. 1. This interactive animation of World Bank data was designed with only 20 lines of Animint code (<http://bit.ly/1rjwX4r>). **Top text:** Five geometric elements span two linked plots: clicking a `clickSelects` element changes the currently selected year or country, and updates the corresponding `showSelected` elements. **Left plot:** a multiple time series from 1960 to 2012 of the life expectancy of 205 countries, with bold lines showing the selected countries and a vertical grey tallrect showing the selected year. **Right plot:** a scatterplot of life expectancy versus fertility rate of all countries. The two text elements show the current selection: year=1975 and country={United States, Vietnam}

WorldBank visualization. D3 uses a data-bind operation followed by several accessor functions, whereas Animint uses a shorter syntax involving a simple `aes` mapping of data variables to geometric attributes. Importantly, Animint is able to express interactivity using the simple `clickSelects` and `showSelected` keywords, whereas D3 would require much more code involving handler functions for mouse click events.

There are many other libraries which can generate web plots with limited interactivity, but which are unable to produce interactive animations. For example, the `rCharts` R package is a wrapper around several JavaScript libraries [Vaidyanathan, 2013]. In contrast to Animint, `rCharts` does not include a DSL for describing interactions between multiple plots. Thus `rCharts` can produce several simple visualizations with only 1 line of R code, but it can not produce multi-plot interactive animations. In contrast, Animint’s high-level DSL can be used to specify more complex visualizations using only tens of lines of R code.

Finally, Dimensional Charting (DC) is another JavaScript library which can produce interactive visualizations consisting of several linked plots [DC.js Team, 2014]. Like Animint, DC builds on top of D3. Unlike Animint, DC does not use the grammar of graphics, so DC designers are limited to plot types that are pre-defined by the library developers.

2.3 Animated graphics libraries

One way to achieve animation in an iterative programming syntax is by using a `for` statement to loop over the time variable. This is the main idea of the animation package [Xie, 2013]. The main difference between this system and Animint is that the only interaction possible with animation is rewinding and fast-forwarding through the animation frames.

Another way to produce an animated scatterplot of the World Bank data is by using a Google motion chart, available in R through the `googleVis` package [Gesmann and de Castillo, 2011]. The main limitation of this system is that it can only produce a few pre-defined plot types, only one of which can be viewed at any time.

2.4 Libraries based on the grammar of graphics

In this section we discuss the differences between Animint and several other high-level DSLs for data visualization based on the grammar of graphics [Wilkinson et al., 2006].

Animint extends the declarative DSL of `ggplot2` [Wickham, 2009, 2010]. Strictly speaking, `ggplot2` is for non-interactive

and non-animated visualizations. In this paper, we propose the `clickSelects` and `showSelected` aesthetics for `ggplot2`, which extend it to accommodate interactive, animated graphics.

Another R package that uses the grammar of graphics to define interactive graphics is called `ggvis` [RStudio, 2014]. Unlike Animint, `ggvis` does not directly extend `ggplot2`, but provides a new implementation of some of the same ideas about the grammar of graphics. There are many other differences between `ggvis` and Animint. The main difference is that `ggvis` can be used to define single plots, but Animint can be used to define multiple plots which are linked using the `clickSelects` and `showSelected` aesthetics. Another difference is that sharing a `ggvis` visualization requires a web server that runs R and shiny, but Animint uses static files and client-side JavaScript so it does not require any special web server software. The two packages also have different interactive features: `ggvis` uses sliders, checkboxes, and other HTML form elements, whereas Animint users can directly click the SVG elements that are used to visualize the data. For example, in a `ggvis` of the WorldBank data, it is natural to use a slider to select the displayed year. In contrast, in Animint we used a multiple time series plot where the year can be selected by directly clicking the data values on the plot (Figure 1). The Animint plot is thus easier for the user since it has less articulatory indirection [Hutchins et al., 1985], and less spatial offset [Beaudouin-Lafon, 2000].

Like Animint, Vega is a declarative DSL that builds on top of D3 [Trifacta, 2014]. Vega is unable to express all plots that can be made using pure D3, but provides a JSON file format capable of defining many common plots (Figure 2). As discussed in Section 3, Animint also internally uses a JSON file to store meta-data about an interactive animation. The main difference is that Vega does not support interactions and animation that show and hide data subsets across multiple linked plots. Finally, Satyanarayan et al. [2014] proposed some declarative Vega extensions “critical for developing interactive data visualizations.” Those extensions are defined at a lower level of abstraction than the `clickSelects` and `showSelected` keywords of Animint.

2.5 Other systems

The system described by Cleveland [1993], Becker and Cleveland [1987] involves interaction via a rectangular brush that can select several data points. Its enhanced linking feature highlights the selected data points across several plots, which is similar to the selection functionality implemented in Animint. There main important difference is that Animint supports a brush for every

D3 (Javascript code)	Vega (JSON file)	Animint (R code)
<pre>svg.selectAll("circle") .data(one_year) .enter().append("circle") .attr("cx", function(d) { return x_scale(d.fertility_rate); }).attr("cy", function(d) { return y_scale(d.life_expectancy); }).attr("r", function(d) { return size_scale(d.population); }).style("fill", function(d) { return color_scale(d.region); })</pre>	<pre>{ "marks": [{ "type": "symbol", "from": { "data": "this_year" }, "properties": { "enter": { "x": { "field": "fertility_rate" }, "y": { "field": "life_expectancy" }, "size": { "field": "population" }, "fill": { "field": "region" } } } }] }</pre>	<pre>geom_point(aes(x=fertility.rate, y=life.expectancy, color=region, size=population, clickSelects=country, showSelected=year), data=WorldBank)</pre>

Fig. 2. Comparison of code used to define points on a scatterplot. Note that the Animint code is shorter and simpler than the D3 and Vega code. Also, Animint implements the proposed `clickSelects` and `showSelected` aesthetics (red), but D3 and Vega do not.

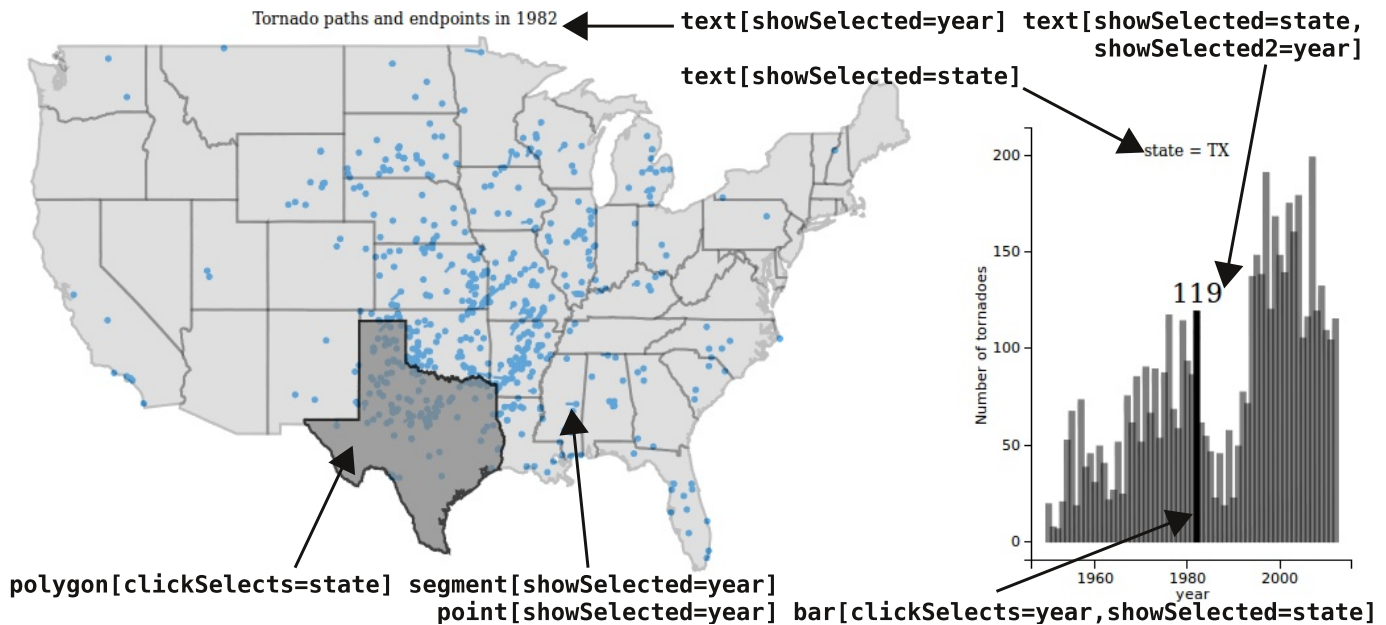


Fig. 3. Interactive animation of tornadoes recorded from 1950 to 2012 in the United States (<http://bit.ly/1hWvYo0>). **Left:** map of the lower 48 United States with tornado paths in 1982. The text shows the selected year, and clicking the map changes the selected state, currently Texas. **Right:** time series of tornado counts in Texas. Clicking a bar changes the selected year, and the text shows selected state and the number of tornadoes recorded there in that year (119 tornadoes in Texas in 1982).

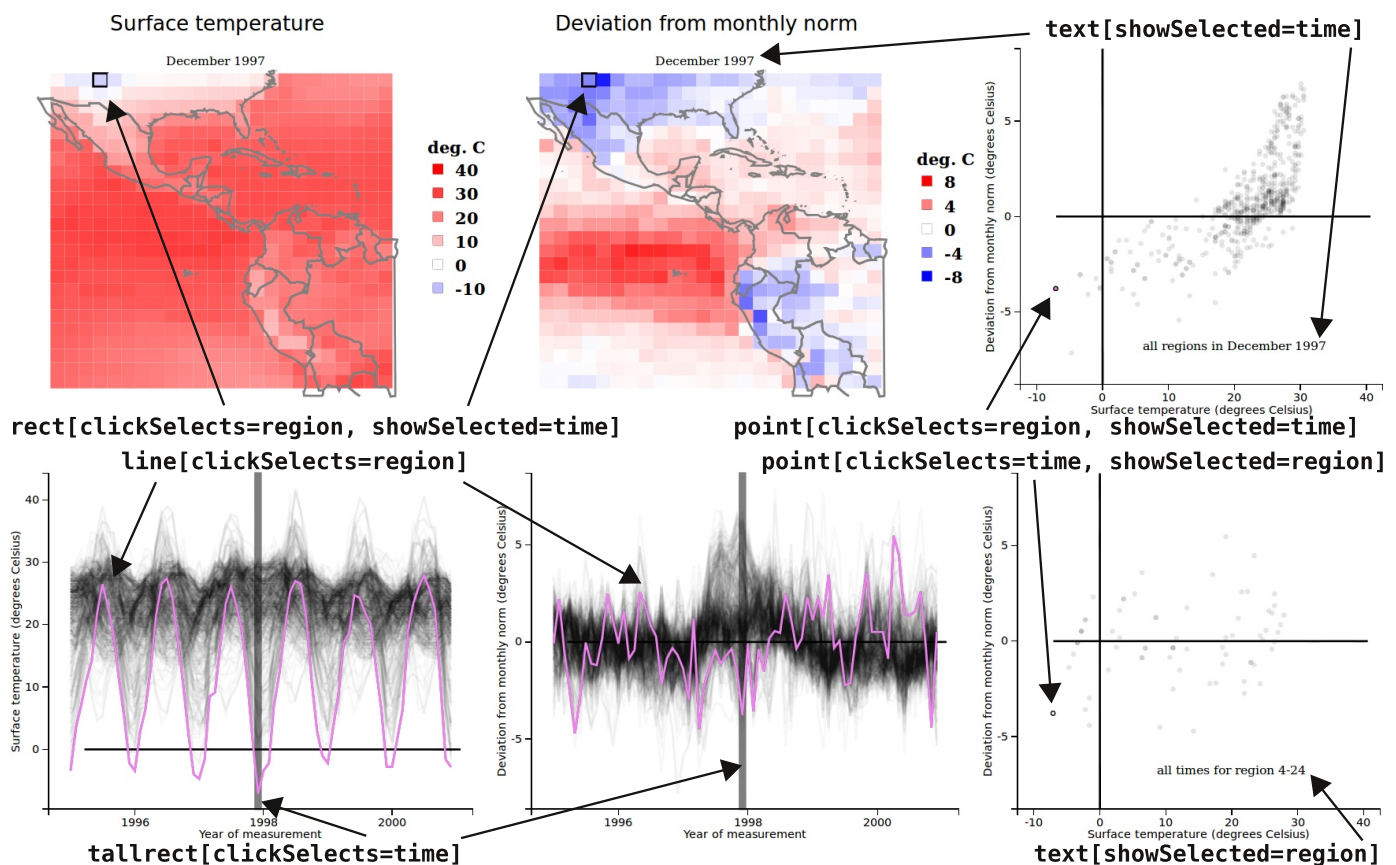


Fig. 4. Visualization containing 6 linked, interactive, animated plots of Central American climate data (<http://bit.ly/QcUrhN>). **Top:** for the selected time (December 1997), maps displaying the spatial distribution of two temperature variables, and a scatterplot of these two variables. The selected region is displayed with a black outline, and can be changed by clicking a rect on the map or a point on the scatterplot. **Bottom:** time series of the two temperature variables with the selected region shown in violet, and a scatterplot of all times for that region. The selected time can be changed by clicking a background tallrect on a time series or a point on the scatterplot. The selected region can be changed by clicking a line on a time series.

selection variable, but Cleveland's system only supports a single brush.

Another system with interactive facilities similar to Animint is Tableau [2014]. Like Animint, it is based on a declarative DSL. The Tableau DSL is called VizQL, which is based on another DSL called Polaris [Stolte and Hanrahan, 2002]. TODO: CARSON DISCUSS TABLEAU.

Finally, implementations of the grammar of graphics system of Wilkinson et al. [2006] exist as the Visualization Markup Language (ViZml) and the Graphics Production Language (GPL) in IBM SPSS Statistics software, but do not support interactive animations.

In conclusion, other systems that implement the grammar of graphics are limited to non-interactive plots, and other systems that implement interactive animations do not exploit the powerful grammar of graphics. Animint is the first system that implements both.

3 THE ANIMINT SYSTEM

In this section we explain the main ideas of how Animint can be used for interactive animations. We first explain the DSL that a designer uses to specify an interactive plot, then explain the user interface for selecting data subsets. Finally, we discuss some implementation details of the compiler and renderer.

3.1 The Animint grammar for interactive animations

The main idea of Animint is that a large class of interactive plots can be specified using just two interactive keywords: `clickSelects` and `showSelected`. For example, after clicking a `clickSelects=year` element, the plot is updated to show only the `showSelected=year` elements corresponding to the selected year. These two simple keywords form the basis of the Animint DSL which makes it easy for designers to translate ideas into code and then visualizations.

The first step in the design of any data visualization is usually to make a sketch of the desired interactive plot on paper or a whiteboard. In any plot a designer would sketch the axes, legends, a few geometric elements, and note which variables will be shown in the linked plots. An Animint designer needs only add the `clickSelects` and `showSelected` mappings for each geometric element, as shown in Figures 1, 3, and 4. These notes can be directly translated to ggplot2 aesthetics in R code, as explained below.

Animint implements the `clickSelects` and `showSelected` interactive keywords as aesthetics in ggplot2. An aesthetic is a declarative mapping from a data variable to a visual property of a geometric plot element. Some standard ggplot2 aesthetics are

<code>x</code>	horizontal position,
<code>y</code>	vertical position,
<code>color</code>	color or fill,
<code>size</code>	point or line thickness.

For example, consider the following R code which defines the points in the World Bank data scatterplot on the right of Figure 1:

```
countryPoints <- geom_point(
  aes(x=fertility.rate, y=life.expectancy,
      color=region, size=population,
      clickSelects=country,
```

```
      showSelected=year),
  data=WorldBank)
```

The code defines a `geom_point`, which means to create a point for every row in the `WorldBank` data table. The visual characteristics of each point are defined by the values of the corresponding data: the (x,y) position encodes fertility rate and life expectancy, point color encodes the region, and point size encodes the population. Note that scales are automatically constructed for the x and y axes, and legends are automatically constructed for color and size. The interactivity is also defined as a simple variable mapping: `clickSelects=country` means that clicking a point changes the selected country, and `showSelected=year` means to only plot the points for the selected year.

In order to remind the plot user which subset of data are selected, we will draw a text label with the selected year and country. First, we create the year labels using

```
yearText <- geom_text(
  aes(label=sprintf("year = %d", year),
      showSelected=year),
  x=5, y=80,
  data=years)
```

Note that `data=years` specifies another data table, with 1 row for each year. Animint does not require that linked plots originate from the same data table; the `clickSelects` and `showSelected` aesthetics will work as long as the different data tables have common variable names (e.g. the `year` variable is present in both `WorldBank` and `years`). The `label` aesthetic is used to define the text from the `year` variable, and only the selected year is shown due to the `showSelected=year` aesthetic. Finally, note that since the `label` x and y positions are constant, they are not defined as aesthetics.

We can also add another label to show the selected country:

```
countryText <- geom_text(
  aes(x=fertility.rate, y=life.expectancy,
      label=country,
      showSelected=country,
      showSelected2=year),
  data=WorldBank)
```

Note that Animint allows any number of `showSelected` aesthetics. In this example, `showSelected=country` combined with `showSelected2=year` means to only show the subset of labels corresponding to both the selected year and country. Since there is only 1 row for each (country,year) combination in the `WorldBank` data, this has the effect of drawing the selected country's label at the location of the selected year.

Having defined these 3 geometric elements, we combine them in a single ggplot:

```
scatterPlot <- ggplot()+
  countryPoints+
  yearText+
  countryText
```

This completes the definition of the scatterplot. Now, we discuss the time series on the left of Figure 1. First, the tallrects in the background are used to select the year:

```
yearReacts <- geom_tallrect(
  aes(xmin=year-1/2, xmax=year+1/2,
      clickSelects=year),
  data=years, alpha=1/2)
```

The tallrect is an Animint extension useful for selecting the variable which is plotted on the x axis, such as `year` in this example. The tallrect plots a rectangle for every row of the `years` data table. The rectangle spans the entire y region, and the `xmin` and `xmax` aesthetics define the left and right limits. Since the designer specified `clickSelects=year`, users can click on a tallrect to select a year.

To create the time series plot, we combine the tallrects above with lines. To declare that clicking a line should change the selected country, we use the `clickSelects=country` aesthetic:

```
timeSeries <- ggplot()+
  yearReacts+
  geom_line(
    aes(x=year, y=life.expectancy,
        group=country, color=region,
        clickSelects=country),
    data=WorldBank, size=3, alpha=3/5)
```

Note that both the `timeSeries` and `scatterPlot` objects are valid ggplots. However, plotting them using the standard `ggplot2` library will show a non-interactive plot with all geometric elements, including data for all years and countries. To plot them with Animint, we define a list of ggplots and options, then call the `animint2dir` compiler:

```
viz <-
  list(scatterPlot=scatterPlot,
       timeSeries=timeSeries,
       time=list(variable="year", ms=3000),
       duration=list(year=1000))
animint2dir(viz, out.dir="WorldBank")
```

The `time` option specifies that in absence of user interaction,

we want the plots to animate over time, progressing at a rate of one year every 3 seconds. We also use the `duration` option to specify a smooth transition over 1 second for the year variable. The `animint2dir` compiler saves some data files and a web page in the `WorldBank` directory, then it opens the interactive plot in a web browser.

3.2 User interaction

Beaudouin-Lafon [2000] discussed the advantages of direct manipulation in graphical user interfaces, and our Animint system follows principal 2 “Physical actions on objects vs. complex syntax” such as dialog boxes. In particular, the user can update the selection by clicking on the data objects themselves. This contrasts other systems which use menus and widgets, and thus suffer from less articulatory directness [Hutchins et al., 1985].

For single selection variables such as year, clicking sets the value of the corresponding selection variable. For multiple selection variables such as country, clicking adds or removes values from the set of selected values.

3.3 Implementation details

As shown in Figure 6, the Animint system is implemented in 2 parts: the compiler and the renderer.

The compiler is implemented in about 1500 lines of R code that converts a list of ggplots and options to a comma-separated values (CSV) file database and a JSON plot meta-data file. The compiler scans the aesthetics in all of the ggplots to determine how many selection variables are present, and which plots to update after a selection variable is clicked. It uses `ggplot2` to automatically calculate the axes scales, legends, and labels. It outputs this information to the JSON plot meta-data file. It also uses `ggplot2` to convert data variables (e.g. life expectancy and region) to visual properties (e.g. y position and color). The data are separated into several CSV files, so for large data sets the web browser only needs to download the subset of data required to render the current selection [Liu et al., 2013]. Finally, the rendering engine (`index.html`, `d3.v3.js`, and `animint.js` files) is copied to the plot directory. Since the compiled plot is just a

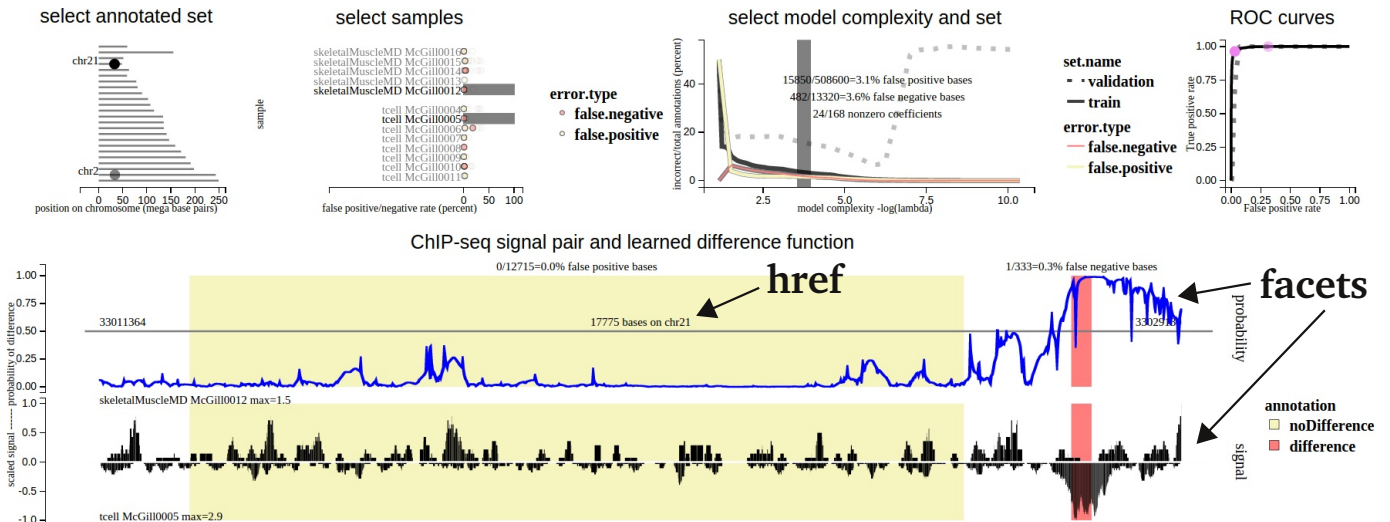


Fig. 5. Visualization with 4 selection variables used to navigate through 1,292,464 rows of data in 5 linked interactive plots (<http://bit.ly/1pVZZaS>). The bottom plot shows a facets plot with aligned x axes, used to emphasize that the blue probability function is defined at the same positions as the black signals below. It also contains an href tag, which opens a new genome browser web page zoomed to the same region as the selected data.

directory of files, the designer can easily upload interactive plots to the web for sharing with users.

The `animint.js` renderer is implemented in about 1500 lines of JavaScript/D3 code that renders the CSV and JSON plot data files as SVG in a web browser. Importantly, animation is achieved by using the JavaScript `setInterval` function, which updates the `time` selection variable every few seconds.

4 RESULTS AND COMPARISON STUDY ON WORLD BANK DATA

To show the advantages that the Animint grammar brings for creating interactive and animated data visualizations, we implemented the World Bank visualization of Figure 1 using two other R packages (Table 1). The main result of our comparison is that Animint requires significantly fewer lines of code, and produces interactive plots with more articulatory directness [Hutchins et al., 1985]. Note that it is possible to implement the World Bank visualization in pure D3 (Figure 2), but would require significantly more code.

4.1 R package animation

We designed a version of the WorldBank visualization with limited interactivity, using 38 lines of R code and the `animation` package (<http://bit.ly/1hnUnkE>). The main idea behind this approach is to use an imperative programming style with for loops to create a static PNG image for each year of the data, and then show these images in sequence. The main drawback to this approach is that the resulting plot is only interactive with respect to the year variable. In other words, the designer must select some countries to emphasize, and the user can not change that selection. Another drawback is that R package `animation` does not support smooth transitions between animation frames. In contrast, using only 20 lines of the Animint DSL, the Animint package achieves smooth transitions and interaction with respect to both year and country variables.

4.2 Client-server systems like ggvis/shiny

We designed another version of the World Bank data visualization in 84 lines of R code, using the `ggvis` graphics library combined

TABLE 1

Implementation complexity and features of the World Bank data visualization using several libraries that can create interactive animations. For each library we show the number of lines of code (LOC), the on-screen objects that can be clicked, the number of interaction variables, and URL of the interactive version.

library	LOC	click on	interaction vars	http://bit.ly/
animint	20	plotted data	several	1rjwX4r
animation	38	play/pause	1 = time	1hnUnkE
ggvis/shiny	84	widgets	several	1diUYsg

with the recommended shiny web server package [RStudio, 2013, 2014] (<http://bit.ly/1diUYsg>). Showing and hiding data subsets was accomplished by clicking on a slider for year and a menu for country, not by clicking on the plot elements. In contrast, we designed Figure 1 using only 20 lines of R code with the Animint package. The reason why the implementation is significantly simpler using Animint is because its DSL is designed specifically for this type of interactive animation.

Another difference is the amount of work required to deploy or share a visualization. A compiled Animint visualization consists of static CSV, JSON, HTML, and JavaScript files which can be easily served with any web server. In contrast, `ggvis+shiny` requires a web server with R installed, significantly complicating deployment to the web.

There are also inherent speed tradeoffs to using a client-server plotting system like `ggvis+shiny` rather than an entirely web client/JavaScript-based system like Animint. There is one main difference between these two types of systems that affects responsiveness of a web-based interactive plotting system: client-server communication overhead. All the Animint JavaScript plot rendering code is executed in the web browser, whereas `ggvis` executes some computations on the server. This means that after a mouse click, `ggvis` can not update a plot immediately, but instead must wait for the server to respond with the plot data.

We quantified speed differences between the two systems by timing web page loading using DevTools in the Chromium web browser Version 33.0.1750.152 Ubuntu 12.04 (256984). We also used `getTime()` in JavaScript to record timings for interactive plot updates (on a desktop computer with a 2.8GHz Intel Core i7 CPU). Using `ggvis` with a local web server and the World Bank data resulted in a web page that loaded quickly (about 1.4s), but updated the plot with a noticeable lag after each mouse click (500–1000ms). Note that since we used a local web server, these times represent the overhead of the web server system, and would be larger with a remote web server.

When we used Animint to make the World Bank data visualization, the compilation from R objects to 2.1MB of uncompressed CSV data files took 2.3s. Using a local web server, the Animint JavaScript rendered the plot very quickly (100–200ms). We also observed very fast plot updates after mouse clicks in Animint: 20–30ms response times for selecting the year, and 60–70ms response times for selecting the country.

The conclusion of our speed comparison is that the overhead of communicating with a web server results in significant slowdowns for interactive animations. It is clear that for quick response times, it is preferable to use an entirely JavaScript-based system like Animint.

In contrast, a web server system like `ggvis+shiny` would be more appropriate for performing arbitrary calculations in R/C code

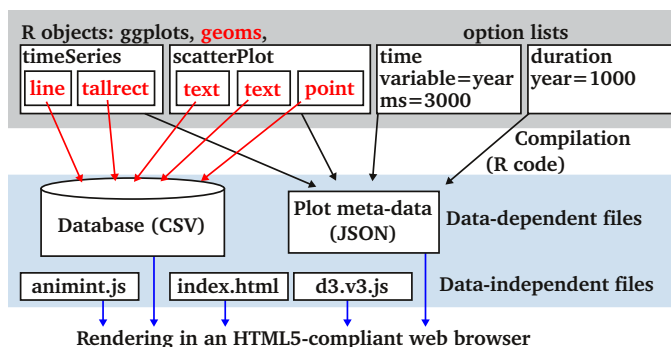


Fig. 6. Schematic explanation of compilation and rendering the World Bank visualization shown in Figure 1. **Top**: the interactive animation is a list of 4 R objects: 2 ggplots and 2 option lists. **Center**: Animint R code compiles data in ggplot geoms to a database of CSV files (→). It also compiles plot meta-data including ggplot aesthetics, animation time options, and transition duration options to a JSON meta-data file (→). **Bottom**: those data-dependent compiled files are combined with data-independent JavaScript and HTML files which render the interactive animation in a web browser (→).

on the server, in response to user inputs, and then sending the result across the network for plotting in the user’s web browser. This power is not always necessary for interactive animations, since the only operation needed is showing precomputed data subsets. However, the web server system would certainly be preferable when there are many more data subsets than could ever be precomputed. In that case, the web server would only compute the subsets that the user interactively specifies.

5 EXAMPLE APPLICATIONS

In this section we discuss the range of examples that we have designed with Animint. Table 2 shows several characteristics of eleven interactive visualizations that we have designed using Animint.

We quantified the implementation difficulty of the Animint examples using lines of R code, including data processing but not including comments (80 characters max per line). We counted the number of plots and variables shown to quantify the amount of information conveyed by the visualization. Using only 17 lines of code, we designed a simple visualization that shows 2 linked plots of 4 variables in the worldPop data set. In contrast, the most complex visualization required 229 lines of code, and it shows 44 variables across 5 linked plots (Figure 5). All of the visualizations that we designed involved at least 2 interaction variables (e.g. year and country in Figure 1) and 2 plots. Indeed, Animint is most appropriate for interactive visualizations of multi-variate data that are not easy to view all at once in one plot.

Table 2 also shows Animint system requirements for plots of various sizes. We timed the compilation step in R code (“seconds” column), and measured the size in megabytes of the compiled CSV file database (“MB” column), and found that both increase with the data set size (“rows” column). We also noticed that the time required for the interactive updates and rendering increases with the amount of data displayed at once (“onscreen” column). In particular, the climate data visualization has noticeably slow animations, since it displays about 88980 geometric elements at once. We observed this slowdown across all browsers, which suggested that there is an inherent bottleneck when rendering large interactive plots in web browsers using JavaScript and SVG. Another Animint with a similar amount of total rows is based on the evolution data (<http://bit.ly/O0VTS4>), but since it shows less data onscreen (about 2703 elements), it exhibits faster responses to interactivity and animation.

5.1 Animated examples

Animation is useful mainly for data sets which have a time variable, as in the World Bank data of Figure 1.

Figure 3 shows an interactive animation of tornadoes observed in the United States between 1950 and 2012. At any moment in time, the user can simultaneously view the spatial distribution of tornadoes in the selected year over all states, and see the trend over all years for the selected state. Clicking a state on the map updates the time series bars to show the tornado counts from that state. Clicking a bar on the time series updates the selected year.

Figure 4 shows an interactive animation of climate time series data observed in Central America. Two maps display the spatial distribution of two temperature variables, which are shown over time in corresponding the time series plots below. Scatterplots also show the relationships between the two temperature variables, for the selected time and region. Clicking any of the plots updates all 6 of them. The `clickSelects` and `showSelected` aesthetics make it easy to design this set of 6 linked plots in only 87 lines of code. Animint’s DSL allows for this level of flexibility while using minimal lines of code to define the plots and the relationship between them.

5.2 Non-animated examples

When the data to visualize do not contain a time variable, we have found that Animint is still useful for creating interactive but non-animated plots. In fact, seven of the eleven examples in Table 2 are not animated. For example, these plots are useful to illustrate complex concepts such as a change point detection model in the breakpoints data (<http://bit.ly/1gGYFIV>). The user can explore different model parameters and data sets since these are encoded as Animint interaction variables.

Another non-animated example is Figure 5, which was used to explain a complex machine learning model for predicting differences between samples. The interactive visualization allows the user to show how the predictions change as a function of the model complexity parameter, in several train and test samples. It also uses facets, a feature from ggplot2 that allows multi-panel plots with aligned axes. Finally, it includes hyperlinks which open related web pages in new windows.

Overall, we have found that Animint is useful for exploring relationships in many different kinds of multivariate data. By using `clickSelects` and `showSelected`, it is easy to design interactive plots that reveal patterns in complex data.

	lines of R code	seconds	MB	rows	onscreen	variables	interactive	plots	animated?	Fig
worldPop	17	0.2	0.1	924	624	4	2	2	yes	
WorldBank	20	2.3	2.1	34132	11611	6	2	2	yes	1
evolution	25	21.6	12.0	240600	2703	5	2	2	yes	
change	36	2.8	2.5	36238	25607	12	2	3	no	
tornado	39	1.7	6.1	103691	16642	11	2	2	no	3
prior	54	0.7	0.2	1960	142	12	3	4	no	
compare	66	10.7	7.9	133958	2140	20	2	5	no	
breakpoints	68	0.5	0.3	4242	667	13	2	3	no	
climate	84	12.8	19.7	253856	88980	15	2	6	yes	4
scaffolds	110	56.3	78.5	618740	9051	30	3	3	no	
ChIPseq	229	29.9	78.3	1292464	1156	44	4	5	no	5

TABLE 2

Characteristics of eleven interactive visualizations designed with Animint. From left to right, we show the data set name, the lines of R code including data processing but not including comments (80 characters max per line), the amount of time it takes to compile the visualization (seconds), the total size of the uncompressed CSV files in megabytes (MB), the total number of data points (rows), the median number of data points shown at once (onscreen), the number of data columns visualized (variables), the number of `clickSelects/showSelected` variables (interactive), the number of linked panels (plots), if the plot is animated, and the corresponding Figure number in this paper (Fig).

6 USER FEEDBACK AND OBSERVATIONS

By working with researchers in several fields of research, we have created a wide variety of interactive visualizations using Animint. Typically, the researchers have a complex data set that they wish to visualize, but they do not have the expertise or time to create an interactive data visualization. The Animint DSL made it easy to collaborate with the various domain experts, who were able to provide us with annotated sketches of the desired plots, which we then translated to R code. In this section we share comments and constructive criticism that we have obtained from our users.

R users have found that Animint is easy to learn. One statistics Ph.D. student writes, “animint is a fantastic framework for creating interactive graphics for someone familiar with R and ggplot2’s grammar of graphics implementation. The API is very intuitive and allows one to quickly bring their static graphics to life in a way that facilitates exploratory data analysis.”

For the prior data visualization (<http://bit.ly/1peIT7t>), the Animint user is a machine learning researcher who developed an algorithm and applied it to 4 benchmark data sets. He wanted to explore how his algorithm performed, in comparison to a baseline learning algorithm. He appreciated the intuition about his algorithm’s performance that he learned from the interactive plots: “Interactive plotting allows us to explore all relationships of our high-dimensional dataset and gives us an intuitive understanding of the performance of our proposed algorithm. An intuitive understanding of the results is important since it shows under which conditions our proposed method works well and provides avenues for further research.”

Another user from a machine learning background found the interactive plots useful for presenting his work: “the ‘regularization path’ is a difficult concept to demonstrate in my research. The [Animint <http://bit.ly/1gVb8To>] helped greatly by rendering an interactive plot of regularization path, likelihood, and graph at the same time and illustrating their connections. It also reveals an interesting phenomenon that maximizing the testing likelihood actually gives many false positives.”

In another application, the Animint user was a genomics researcher: “viewing and exploring my complex intestinal microbiome dataset in [Animint] allowed me to grasp the patterns and relationships between samples at an almost intuitive level. The interactive aspect of it was very helpful for browsing through the dataset.”

Finally, users also appreciated the simple web interface, and the detail that is possible to show in interactive plots, but impossible to show in publications: “... the web interface is simple and easy to use. It also enables us to publish more detailed interactive results on our website to accompany the results presented in publications.”

7 DISCUSSION, LIMITATIONS AND FUTURE WORK

There are several limitations to the current implementation of Animint, which suggest avenues for future work.

Animint implements several linked brushes, one of the hallmarks of interactive visual analysis [Konyha et al., 2012]. However, the main drawback to the current implementation is that a selection is defined as a set of distinct elements (e.g. `year={1991, 1992}`) rather than a logical expression (e.g. `year > 1990`). Importantly, this is a drawback of the current implementation, not the Animint DSL.

Another limitation of the current implementation is that axes and legends are computed at compile-time, and are not updated when the selection changes. Thus it is difficult to use interactivity to visualize different data subsets which have very different ranges of values. For example, it would be awkward to use `showSelected` to show one panel of a scatter plot matrix of several variables. A workaround is shown in Figure 5, which omits the x axis on the bottom plot, since in fact the x values are all normalized to [0,1]. A future implementation of Animint would benefit from scales that update after each click.

Currently, the user is unable to change variable mappings, since these are specified by the designer at compile time. A better interactive system would allow the user to interactively change variable mappings.

Since Animint does not perform any computations other than showing and hiding data subsets, there are some limitations to what can be displayed with multiple selection variables. For instance, in the WorldBank visualization of Figure 1, it would not be feasible to display a single smoothing line computed from all the selected countries. This is because `showSelected=country` means to show one thing for each selected country (not one thing computed based on the set of selected countries). Supporting this kind of interaction would require substantial modifications to the Animint system, including the addition of a web server to perform computations on the multiple selection variable sets.

As discussed in Section 3 and illustrated in Figure 6, the compiler is written in R, and the renderer is written in JavaScript. This means that Animint developers must be proficient in both R and JavaScript. This represents a significant barrier for source code contributions from developers who are proficient with one language but not the other.

Animint designers define interactive animations using only R code, and no knowledge of JavaScript is necessary. This is convenient for users from a statistical background, but presents a barrier for web developers who are more familiar with JavaScript than R. For these web developers, it would be advantageous in the future to implement a compiler and renderer in pure JavaScript, by possibly building `clickSelects` and `showSelected` extensions into Vega [Trifacta, 2014].

The current Animint implementation is limited to two specific types of interactivity: highlighting the selected `clickSelects` element, and showing/hiding `showSelected` elements. In the future, we could implement several other types of interactivity without changing the Animint DSL. Examples include zooming, increase/decrease animation speed, and plot resizing. However, some kinds of interactivity would require additions to the Animint grammar. For example, a `tooltip` aesthetic could be used to indicate the text the user views when they hover the mouse over a data point.

Animint’s performance can be measured using speed, memory, and disk space requirements in the compilation and rendering steps. Although we showed in Section 4 that Animint provides smoother interactivity than client-server systems, future versions of Animint could definitely be made even more efficient and responsive. For example, of the plots in Table 2, the longest compilation step took 56.3 seconds, which may be reduced by optimizing the R code compiler.

This highlights one of the main motivations for using a declarative DSL like Animint: none of the designer’s R code needs to be changed to implement improvements like this. Instead, the Animint developers just need to work on a better compiler and

rendering engine. Indeed, Heer and Bostock [2010] noted that this is one of the main benefits of declarative language design: “By decoupling specification from implementation, developers can implement language optimizations without interfering with the work of designers.”

While several optimizations remain to be implemented, the current Animint library already provides an efficient syntax for the design of interactive, animated data visualizations.

ACKNOWLEDGEMENTS

The authors wish to thank Animint users MC Du Plessis, Song Liu, Nikoleta Juretic, and Eric Audemard who have contributed constructive criticism and helped its development.

REFERENCES

- M. Beaudouin-Lafon. Instrumental interaction: An interaction model for designing post-wimp user interfaces. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, CHI '00, pages 446–453, New York, NY, USA, 2000. ACM. ISBN 1-58113-216-6. doi: 10.1145/332040.332473. URL <http://doi.acm.org/10.1145/332040.332473>.
- R. Becker and W. Cleveland. Brushing scatterplots. *Technometrics*, 29(2):127–142, May 1987.
- M. Bostock, V. Oglevevsky, and J. Heer. D3 data-driven documents. *IEEE Transactions on Visualization and Computer Graphics*, 17(12):2301–2309, December 2011.
- W. S. Cleveland. *Visualizing Data*. Hobart Press, New Jersey, 1993.
- DC.js Team. Dimensional charting JavaScript library, 2014. URL <http://dc-js.github.io/dc.js/>.
- M. Gesmann and D. de Castillo. googleVis: Interface between R and the Google Visualisation API. *The R Journal*, 3(2):40–44, December 2011.
- J. Heer and M. Bostock. Declarative language design for interactive visualization. *IEEE Transactions on Visualization and Computer Graphics*, 16(6):1149–1156, 2010.
- J. Heer and G. Robertson. Animated transitions in statistical data graphics. *IEEE Transaction on Visualization and Computer Graphics*, 13(6):1240–1247, 2007.
- J. Heer, M. Agrawala, and W. Willett. Generalized selection via interactive query relaxation. In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, pages 959–968. ACM, 2008.
- E. L. Hutchins, J. D. Hollan, and D. A. Norman. Direct manipulation interfaces. *Hum.-Comput. Interact.*, 1(4):311–338, Dec. 1985. ISSN 0737-0024. doi: 10.1207/s15327051hci0104_2. URL http://dx.doi.org/10.1207/s15327051hci0104_2.
- Z. Konyha, A. Lež, K. Matković, M. Jelović, and H. Hauser. Interactive visual analysis of families of curves using data aggregation and derivation. In *Proceedings of the 12th International Conference on Knowledge Management and Knowledge Technologies*, page 24. ACM, 2012.
- Z. Liu, B. Jiang, and J. Heer. immens: Real-time visual querying of big data. *Computer Graphics Forum (Proc. EuroVis)*, 32, 2013. URL <http://vis.stanford.edu/papers/immens>.
- RStudio. shiny: easy web applications in R, 2013. URL <http://www.rstudio.com/shiny/>.
- RStudio. ggvis: interactive grammar of graphics for R, Mar 2014. URL <http://ggvis.rstudio.com>.
- A. Satyanarayan, K. Wongsuphasawat, and J. Heer. Declarative interaction design for data visualization. In *ACM User Interface Software & Technology (UIST)*, 2014. URL <http://idl.cs.washington.edu/papers/reactive-vega>.
- B. Shneiderman. The future of interactive systems and the emergence of direct manipulation. *Behaviour & Information Technology*, 1(3):237–256, 1982.
- C. Stolte and P. Hanrahan. Polaris: A system for query, analysis and visualization of multi-dimensional relational databases. *IEEE Transactions on Visualization and Computer Graphics*, 8:52–65, 2002.
- Tableau. VizQL: Visualization Query Language, 2014. URL <http://www.tableausoftware.com/>.
- Trifacta. Vega: a declarative visualization grammar, Mar 2014. URL <http://trifacta.github.io/vega/>.
- R. Vaidyanathan. rCharts: Interactive JS Charts from R, 2013. URL <https://github.com/ramnathv/rCharts>.
- H. Wickham. *ggplot2: elegant graphics for data analysis*. Springer New York, 2009. ISBN 978-0-387-98140-6. URL <http://had.co.nz/ggplot2/book>.
- H. Wickham. A layered grammar of graphics. *Journal of Computational and Graphical Statistics*, 19(1):3–28, 2010.
- L. Wilkinson, D. Wills, D. Rope, A. Norton, and R. Dubbs. *The grammar of graphics*. Springer, 2006.
- W. Willett, J. Heer, and M. Agrawala. Scented widgets: Improving navigation cues with embedded visualizations. *IEEE Transactions on Visualization and Computer Graphics*, 13(6):1129–1136, Nov. 2007. ISSN 1077-2626. doi: 10.1109/TVCG.2007.70589. URL <http://dx.doi.org/10.1109/TVCG.2007.70589>.
- World Bank. World development indicators, 2012. URL <http://data.worldbank.org/data-catalog/world-development-indicators>.
- Y. Xie. animation: An R package for creating animations and demonstrating statistical methods. *Journal of Statistical Software*, 53(1):1–27, 2013. URL <http://www.jstatsoft.org/v53/i01/>.